

*Statistical View on Machine Learning*

---

# CNN – Convolutional Neural Networks

Marta Kałużna  
Szymon Czop

---



---

# Plan of our presentation

---

1. Why using CNN is better than normal Neural Networks?
2. CNN overview. Convolution, ReLU and pooling - what do they mean?
3. Building and training the model - Jupyter coding.



---

# Before we begin - some facts about CNN

---



---

# Before we begin - some facts about CNN

---

- ❖ One of the variants of neural networks.



---

# Before we begin - some facts about CNN

---

- ❖ One of the variants of neural networks.
- ❖ Used heavily in the field of Computer Vision.



---

# Before we begin - some facts about CNN

---

- ❖ One of the variants of neural networks.
- ❖ Used heavily in the field of Computer Vision.
- ❖ Derives its name from the type of hidden layers.



Why using Convolutional Neural Networks  
is better than normal Neural Networks?



---

# 1. Why using CNN is better than normal NN?

---

- ❖ **Reason 1: Images are Big**



---

# 1. Why using CNN is better than normal NN?

---

## ❖ Reason 1: Images are Big

The nice thing about images: pixels are most useful in the context of their neighbors.



---

# 1. Why using CNN is better than normal NN?

---

- ❖ **Reason 1: Images are Big**

The nice thing about images: pixels are most useful in the context of their neighbors.

- ❖ **Reason 2: Positions can change**



---

# 1. Why using CNN is better than normal NN?

---

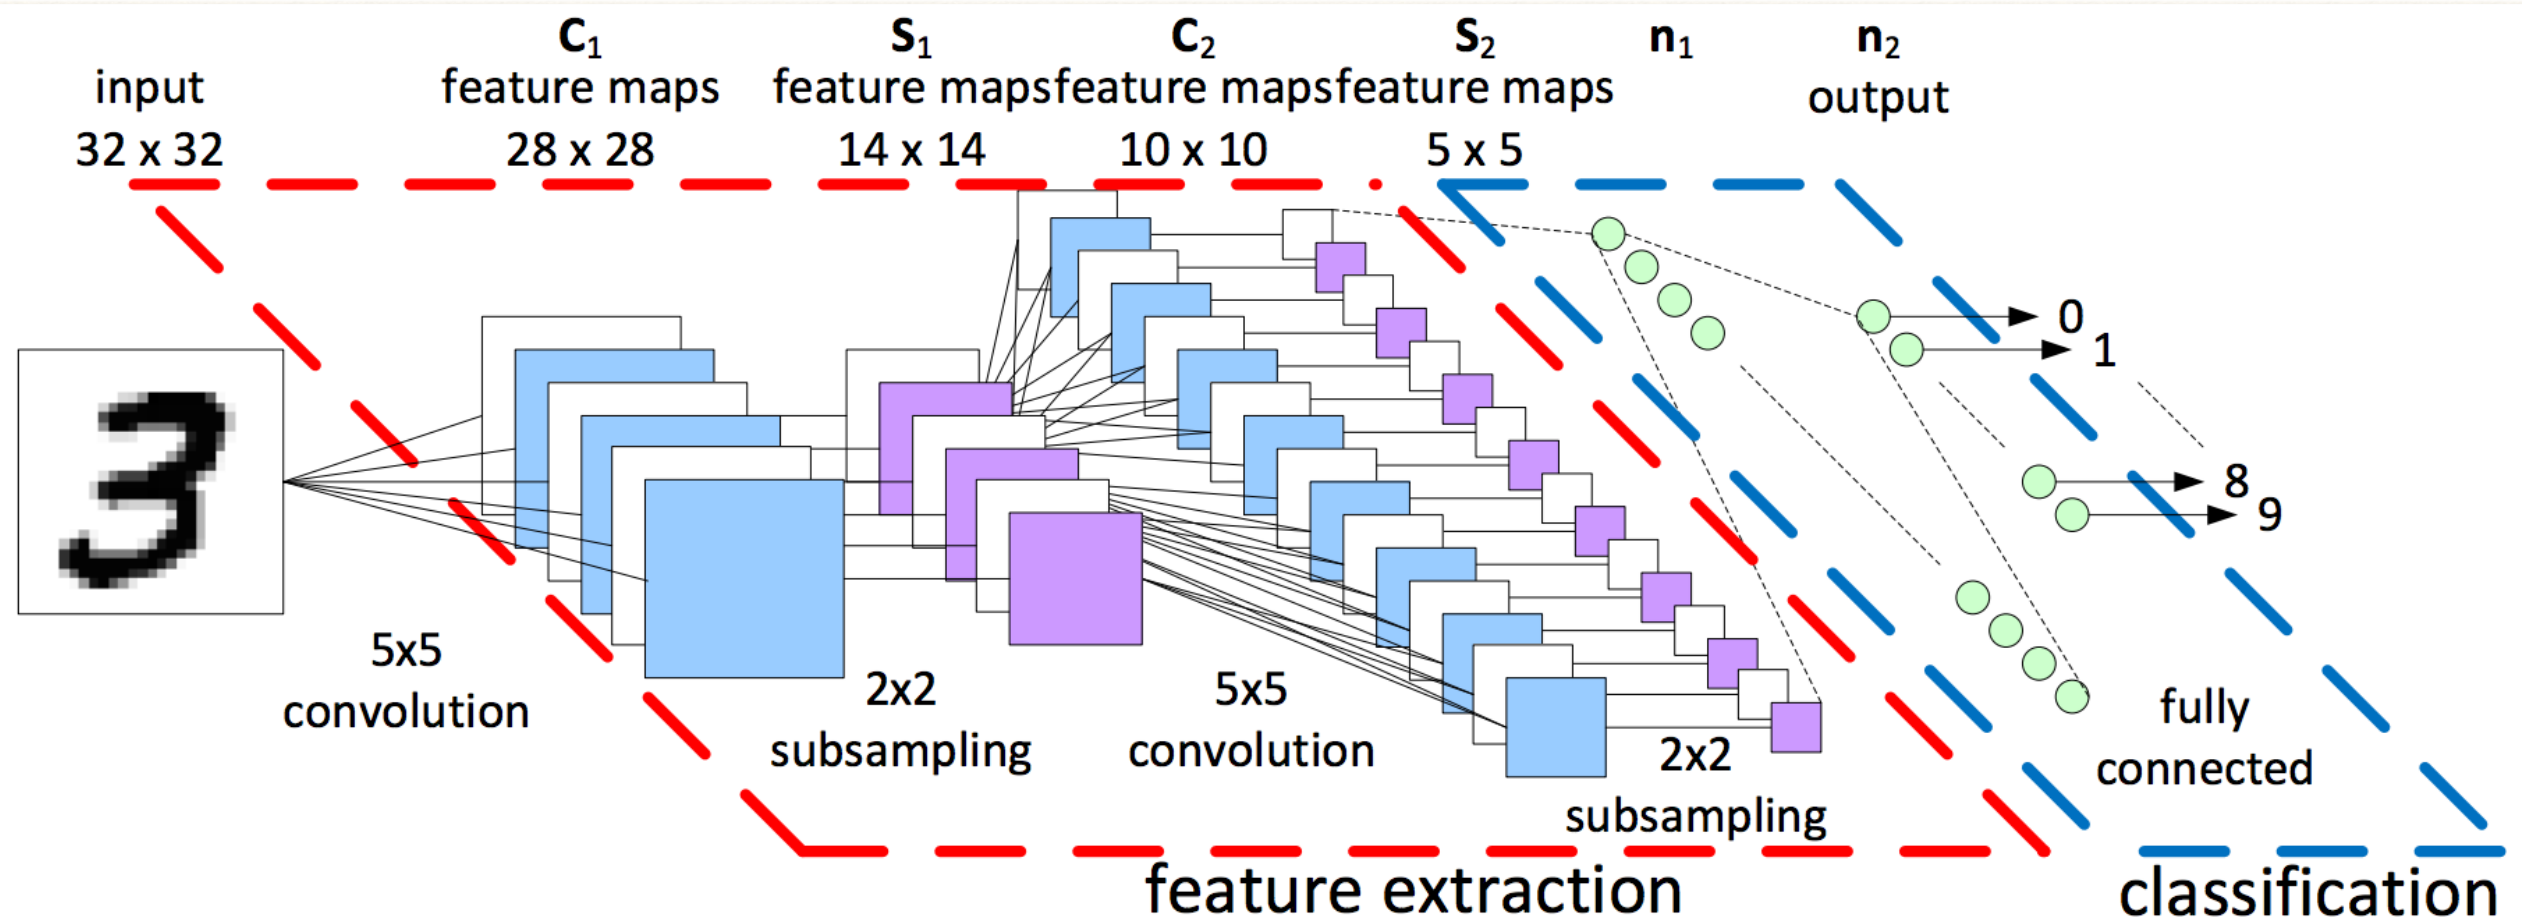
## ❖ Reason 1: Images are Big

The nice thing about images: pixels are most useful in the context of their neighbors.

## ❖ Reason 2: Positions can change

We want to be able to detect a thing regardless of where it appears in the image.





## 2. CNN Overview

Convolutional  
Neural Networks  
architecture



---

# Differences between CNN and NN

---



---

# Differences between CNN and NN

---

- ❖ The layers are organized in 3 dimensions: width, height, depth.



---

# Differences between CNN and NN

---

- ❖ The layers are organized in 3 dimensions: width, height, depth.
- ❖ The neurons in one layer do not connect to all the neurons in the next layer (only to a small region of it).



---

# Differences between CNN and NN

---

- ❖ The layers are organized in 3 dimensions: width, height, depth.
- ❖ The neurons in one layer do not connect to all the neurons in the next layer (only to a small region of it).
- ❖ The final output will be reduced to a single vector of probability scores (organized along the depth dimension).



Convolution, padding, pooling



# Convolution

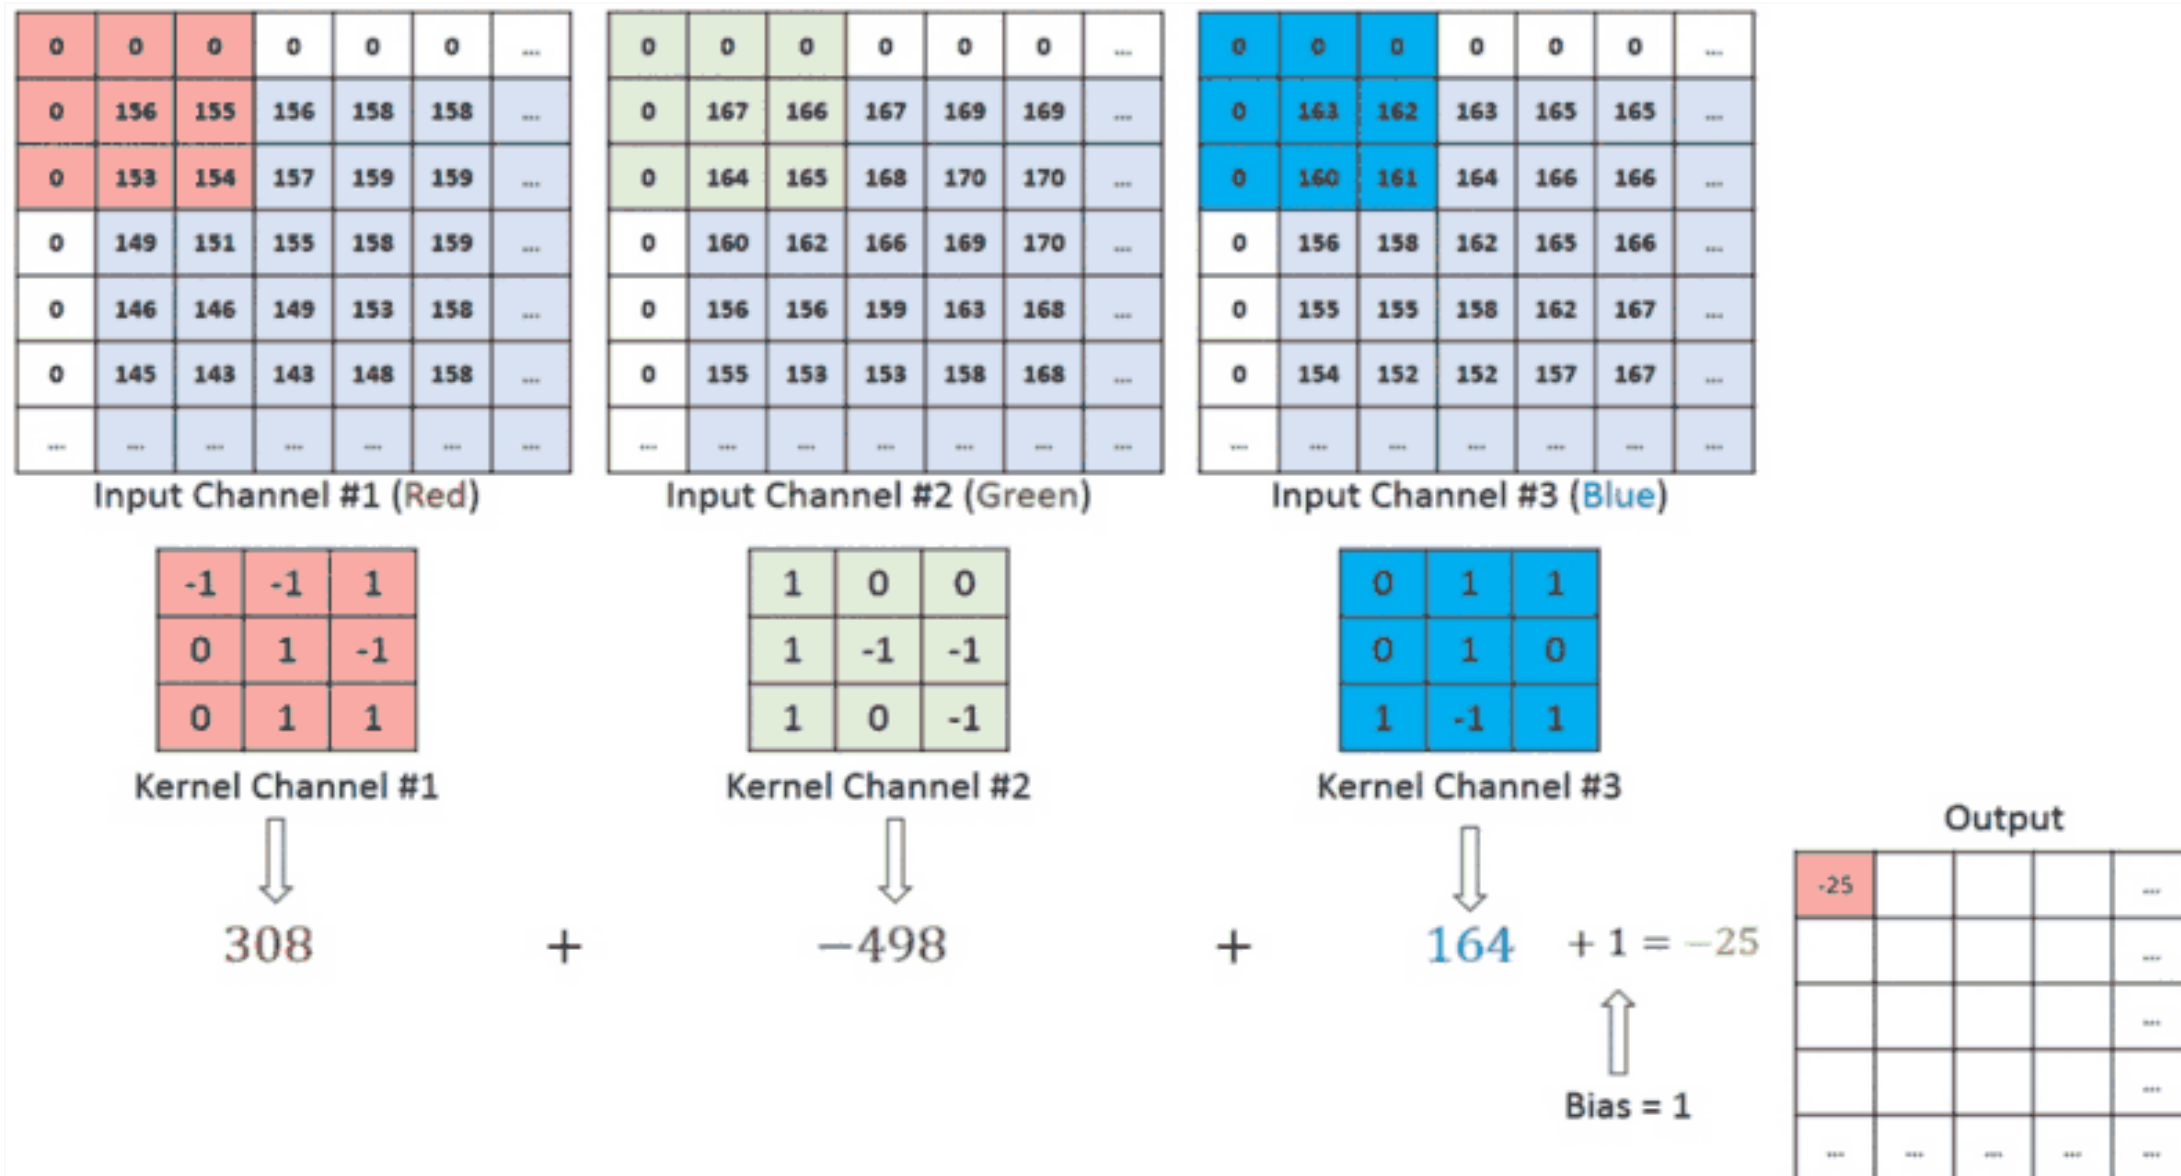
1x1	1x0	1x1	0	0
0x0	1x1	1x0	1	0
0x1	0x0	1x1	1	1
0	0	1	1	0
0	1	1	0	0

4		

1. Overlaying the filter on top of the image at some location.
2. Performing **element-wise multiplication** between the values in the filter and their corresponding values in the image.
3. Summing up all the element-wise products. This sum is the output value for the **destination pixel** in the output image.
4. Repeating for all locations.




# Convolution of images with multiple channels






# How do we connect our filters together?




Convolutional layer with four 3x3 filters on a **black and white image** (just one channel)



Convolutional layer with four 3x3 filters on an **RGB image**. As you can see, the filters are now cubes, and they are applied on the full depth of the image..



# Horizontal vs vertical filter




Image

$$\ast \begin{array}{c} \blacksquare \\ \blacksquare \end{array} =$$

Kernel




Output



Image

$$\ast \begin{array}{cc} \blacksquare & \blacksquare \end{array} =$$






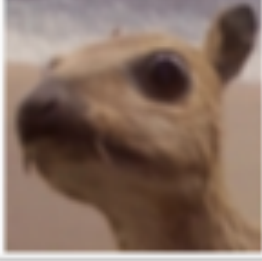
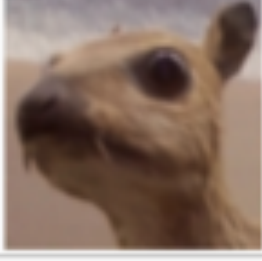
Kernel



Output




Various convolution  
image after applying  
different types of filters

Operation	Filter	Convolved Image
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	
	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	
	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$	
Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	
Box blur (normalized)	$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	
Gaussian blur (approximation)	$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$	



# Padding

- ❖ **Same Padding** - the output image has the same dimensions as the input image (to achieve it, we pad the input image with zeros).





# Padding

0	50	0	29
0	80	31	2
33	90	0	75
0	9	0	95

29	?
?	?

- ❖ **Valid Padding** - the output image is reduced in the dimensionality as compared to the input



# Non Linearity (ReLU)

- ❖ **ReLU** = Rectified Linear Unit
- ❖ The output is  $f(x) = \max(0, x)$ .
- ❖ Converts all of the negative values to 0 and keeps the positive values the same.

1	14	-9	4
-2	-20	10	6
-3	3	11	1
2	54	-2	80



1	14	0	4
0	0	10	6
0	3	11	1
2	54	0	80



---

# Pooling

---

- ❖ Reduces the number of parameters when the images are too large.
- ❖ Why? To decrease the computational power required to process the data through dimensionality reduction.
- ❖ Shortens the training time and controls over-fitting.




# Pooling

- ❖ **Max Pooling** - returns the maximum value from the portion of the image covered by the Kernel.
- ❖ **Average Pooling** - returns the average of all the values from the portion of the image covered by the Kernel.

## Max Pool

2	3	1	9
4	7	3	5
8	2	2	2
1	3	4	5




7	9
8	5

Max-Pool with a 2 by 2 filter and stride 2.

## Average Pool

2	3	1	9
4	7	3	5
8	2	2	2
1	3	4	5

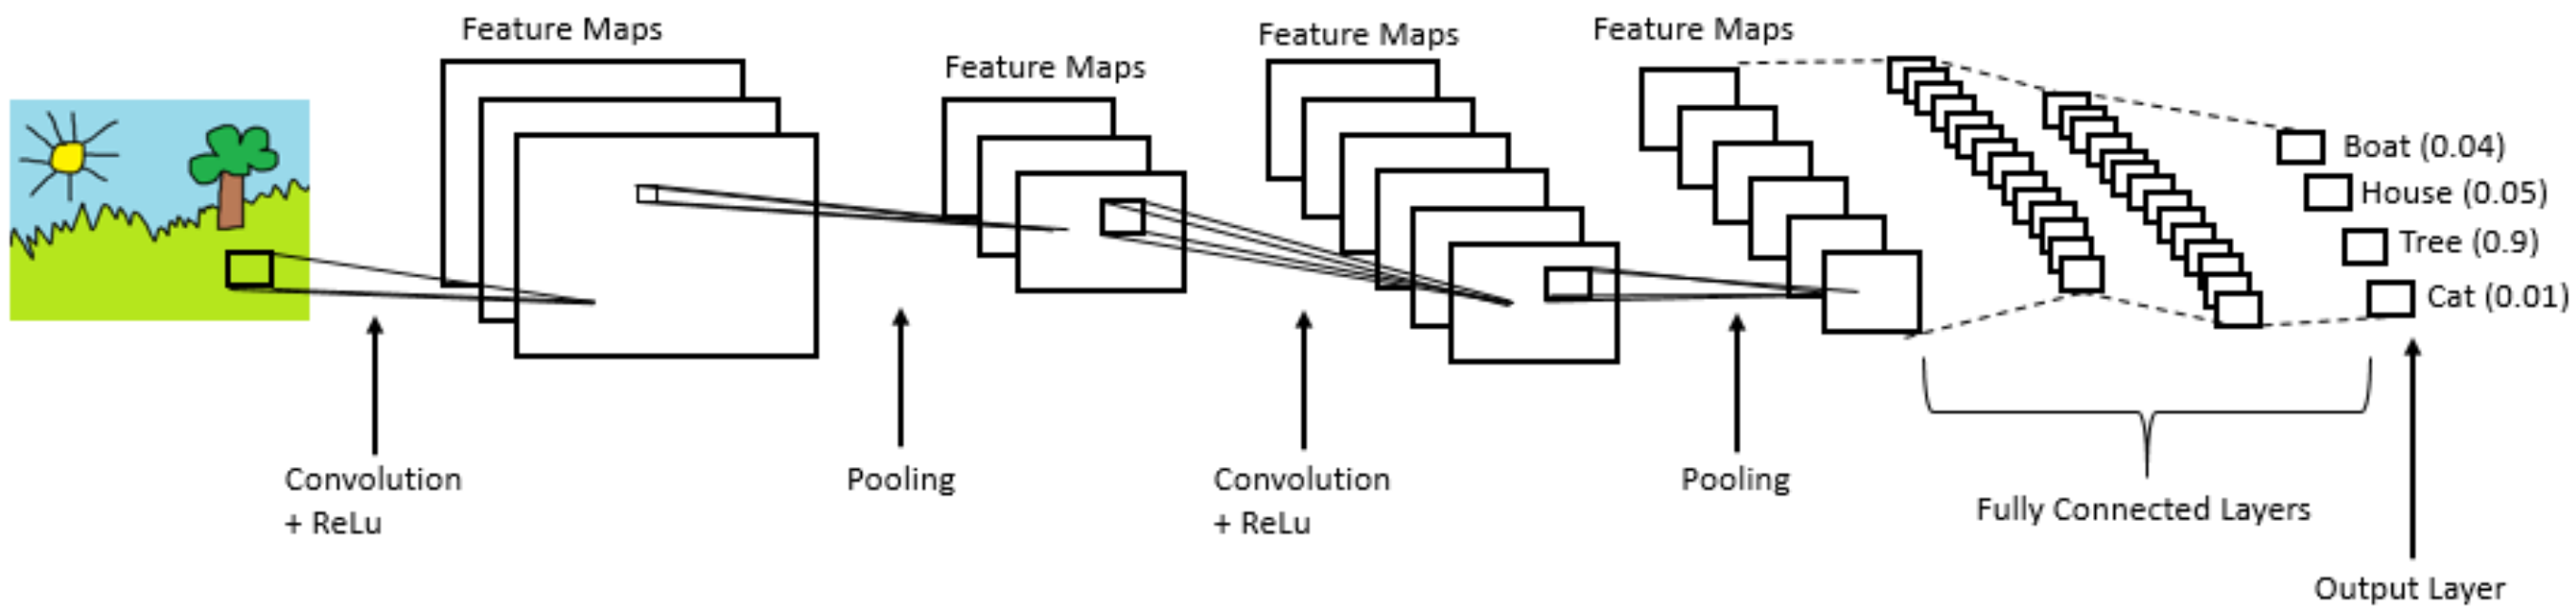


4	4.5
3.25	3.25

Average Pool with a 2 by 2 filter and stride 2.



# Complete CNN architecture






# Backpropagation

## Max Pooling

0	55	0	0
20	0	41	33
0	90	0	0
0	57	0	95



55	41
90	95



## Conv layer

0	7	0
0	80	31
33	14	0

0	0	0
0	1	0
0	0	0



80

$$\begin{aligned} \text{out}(i, j) &= \text{convolve}(\text{image}, \text{filter}) \\ &= \sum_{x=0}^3 \sum_{y=0}^3 \text{image}(i+x, j+y) * \text{filter}(x, y) \\ \frac{\partial \text{out}(i, j)}{\partial \text{filter}(x, y)} &= \text{image}(i+x, j+y) \end{aligned}$$

We can put it all together to find the loss gradient for specific filter weights:

$$\frac{\partial L}{\partial \text{filter}(x, y)} = \sum_i \sum_j \frac{\partial L}{\partial \text{out}(i, j)} * \frac{\partial \text{out}(i, j)}{\partial \text{filter}(x, y)}$$



Building and training the model



Thank you for your attention!